

# Bandit Algorithms

## Learning Under Uncertainty

Antonio Coppola

[ancoppola@unisa.it](mailto:ancoppola@unisa.it)

April 10, 2022

# The Bandit Algorithms

# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell



# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell
- You decide to organize a street market for selling the shoes



# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell
- You decide to organize a street market for selling the shoes
- You post the price 20€ and sell a pair of shoes to any customer that are willing to pay the posted price. Each customer can buy only a pair of shoes



# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell
- You decide to organize a street market for selling the shoes
- You post the price 20€ and sell a pair of shoes to any customer that are willing to pay the posted price. Each customer can buy only a pair of shoes
- Each customer that has a private valuation greater or equal than 20€ buys the shoes



# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell
- You decide to organize a street market for selling the shoes
- You post the price 20€ and sell a pair of shoes to any customer that are willing to pay the posted price. Each customer can buy only a pair of shoes
- Each customer that has a private valuation greater or equal than 20€ buys the shoes



The goal is to maximize the revenue

# Posted Price Auctions

- Suppose that you have a huge amount of shoes to sell
- You decide to organize a street market for selling the shoes
- You post the price 20€ and sell a pair of shoes to any customer that are willing to pay the posted price. Each customer can buy only a pair of shoes
- Each customer that has a private valuation greater or equal than 20€ buys the shoes



The goal is to maximize the revenue

You want sell the shoes at highest possible price,  
but a too high price might make you earn nothing



# Repeated Posted Price Auctions

- Suppose now that you know an important customer that every Monday asks you the price of a pair of shoes
- Each week the customer can change its valuation of the shoes and you can change the price



# Repeated Posted Price Auctions

- Suppose now that you know an important customer that every Monday asks you the price of a pair of shoes
- Each week the customer can change its valuation of the shoes and you can change the price
- Every Monday:
  - 1 The customer has its own valuation



# Repeated Posted Price Auctions

- Suppose now that you know an important customer that every Monday asks you the price of a pair of shoes
- Each week the customer can change its valuation of the shoes and you can change the price
- Every Monday:
  - 1 The customer has its own **valuation**
  - 2 You **post the price** for a pair of shoes



# Repeated Posted Price Auctions

- Suppose now that you know an important customer that every Monday asks you the price of a pair of shoes
- Each week the customer can change its valuation of the shoes and you can change the price
- Every Monday:
  - ① The customer has its own **valuation**
  - ② You **post the price** for a pair of shoes
  - ③ If the **posted price is lower or equal than customer's valuation**, then you sell the pair of shoes at your price



# Repeated Posted Price Auctions

- Suppose now that you know an important customer that every Monday asks you the price of a pair of shoes
- Each week the customer can change its valuation of the shoes and you can change the price
- Every Monday:
  - ① The customer has its own **valuation**
  - ② You **post the price** for a pair of shoes
  - ③ If the **posted price is lower or equal than customer's valuation**, then you sell the pair of shoes at your price



Every Monday you want to post a price to maximize the **cumulated revenue**

# Repeated Posted Price Auctions: Formalization

- For simplicity we consider a two-player game between the seller and the buyer
- At each time step  $t$  the seller wants to sell the good to the buyer
- For each time step  $t \in \{1, 2, \dots, T\}$ :
  - 1 Buyer arrives with (hidden) valuation  $v_t \in [0, P]$
  - 2 Seller sets the price  $\pi_t \in [0, P]$
  - 3 If  $\pi_t \leq v_t$ , then the buyer buys the good and pays  $\pi_t$



The Seller's Goal is to **maximize the cumulated** revenue

$$\sum_{t=1}^T r_t = \sum_{t=1}^T \pi_t \cdot \mathbb{1}\{\pi_t \leq v_t\}$$

# Repeated Posted Price Auctions: A Classic Dilemma

Let's back to our example.

You are the Seller and you can post only two prices  $\{10\text{€}, 20\text{€}\}$  for your pair of shoes. You are running the Repeated Posted Price Auction and you have already posted each price 5 times:

Step	1	2	3	4	5	6	7	8	9	10
Posted Price = 10	10€		10€	0		0				10€
Posted Price = 20		20€			0		0	0	0	

- Cumulated revenue for  $\{\text{Price} = 10\} = 30\text{€}$
- Cumulated revenue for  $\{\text{Price} = 20\} = 20\text{€}$

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?



# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$
- Would you attribute the poor performance of  $\text{Price} = 20\text{€}$  to bad luck and try it a few more times?

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$
- Would you attribute the poor performance of  $\text{Price} = 20\text{€}$  to bad luck and try it a few more times? How many times?

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$  → Exploitation
- Would you attribute the poor performance of  $\text{Price} = 20\text{€}$  to bad luck and try it a few more times? How many times?

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$  → **Exploitation**
- Would you attribute the poor performance of  $\text{Price} = 20\text{€}$  to bad luck and try it a few more times? How many times? → **Exploration**

# Repeated Posted Price Auctions: A Classic Dilemma

You have 10 more trials altogether. What is your strategy?

- Will you keep posting  $\text{Price} = 10\text{€}$  , ignoring  $\text{Price} = 20\text{€}$  → **Exploitation**
- Would you attribute the poor performance of  $\text{Price} = 20\text{€}$  to bad luck and try it a few more times? How many times? → **Exploration**

Finding the right balance between exploration and exploitation is at the heart of all **Bandit** algorithms.

## What are Bandits Algorithms?

## What are Bandits Algorithms?

*"Bandits is a simple but very powerful framework for algorithms that make decisions over time under uncertainty"*

---

Aleksandrs Slivkins. *Introduction to Multi-Armed Bandits*. Foundations and Trends in Machine Learning, 2019



# Some History

- Bandit problems were introduced by William R. Thompson in 1933
- The name comes from the 1950s, when Frederick Mosteller and Robert Bush decided to study how animals learn online
- **Why Bandit?** In first studies, a 'two-armed bandit' was used to model how humans learn online.




- Bandit Algorithms are used by major tech (e.g. Amazon<sup>1</sup>, Zillow<sup>2</sup>) for configuring web interfaces:
  - News Recommendation
  - Dynamic Pricing
  - Ad Placement

---

<sup>1</sup>Daniel N. Hill et al. “An Efficient Bandit Algorithm for Realtime Multivariate Optimization”. In: *CoRR* abs/1810.09558 (2018). arXiv: 1810.09558. URL: <http://arxiv.org/abs/1810.09558>.

<sup>2</sup>Luca Cazzanti. *Contextual, multi-armed bandit performance assessment*. Zillow. 2020. URL: <https://www.zillow.com/tech/multi-armed-bandits/>


## Netflix Artwork Personalization



Ashok Chandrashekar et al. *Artwork Personalization at Netflix*. Netflix. 2017. URL: <https://netflixtechblog.com/artwork-personalization-c589f074ad76>

- Bandit Algorithms are used by major tech (e.g. Amazon, Zillow) for configuring web interfaces:
  - News Recommendation
  - Dynamic Pricing
  - Ad Placement
- Bandit Algorithms can be used to address emerging problems in different disciplines:
  - Auctions
  - Diffusion Mechanisms

# The Language of Bandits



- The bandit problem is a sequential game between a **learner** (or an algorithm) and an **environment**
- The game is played over  $T$  rounds,  $T$  is called the **horizon**
- At each round the learner chooses an **arm**  $a_t$  collects a **reward** (or cost) and receives a **feedback** from the environment.
- With  $K$  arms we refer to the Bandit Problems as  **$K$ -Armed Bandit**. In general, with more than two arms we use the term **Multi-Armed Bandit**
- The learner's goal is to **maximize its total reward over the  $T$  rounds**.

# The Language of Bandits

---

## The Bandits Problem

---

**Require:**  $K$  arms,  $T$  rounds

- 1: **for**  $t \in [T]$  **do**
  - 2:     Algorithm chooses some arm  $a_t \in [K]$
  - 3:     Environment determines reward  $r_t$
  - 4:     Algorithm receives  $r_t$  and a feedback
  - 5: **end for**
-

# Repeated Posted Price Auctions as Bandit Problems

# Repeated Posted Price Auctions as Bandit Problems

- Seller



# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm
- Revenue

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm
- Revenue  $\rightarrow$  Reward

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm
- Revenue  $\rightarrow$  Reward
- Has the good been sold?



# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm
- Revenue  $\rightarrow$  Reward
- Has the good been sold?  $\rightarrow$  Feedback

# Repeated Posted Price Auctions as Bandit Problems

- Seller → Learner
- Buyer → Environment
- Posted Price → Arm
- Revenue → Reward
- Has the good been sold? → Feedback
- Cumulated Revenue

# Repeated Posted Price Auctions as Bandit Problems

- Seller  $\rightarrow$  Learner
- Buyer  $\rightarrow$  Environment
- Posted Price  $\rightarrow$  Arm
- Revenue  $\rightarrow$  Reward
- Has the good been sold?  $\rightarrow$  Feedback
- Cumulated Revenue  $\rightarrow$  Total Reward over the  $T$  rounds

# Repeated Posted Price Auctions as Bandit Problems

---

## Bandit Posted Price Auctions

---

**Require:**  $K$  prices,  $T$  rounds

- 1: **for**  $t \in [T]$  **do**
  - 2:     Buyer has (hidden) valuation  $v_t$
  - 3:     Seller posts a price  $a_t \in [K]$
  - 4:     Reward is  $r_t = a_t$  if  $v_t \geq a_t$ , otherwise  $r_t = 0$
  - 5:     Seller receives  $r_t$ , and knows if the good has been sold (feedback)
  - 6: **end for**
-

# Type of Environments

Two type of environments:

- **Stochastic**: Given the arm chosen by the learner, the environment samples the reward from an (unknown) distribution
- **Deterministic Adversarial**: The environment determines the rewards deterministically, independently from the arms chosen by the learner

# Posted Price with Stochastic Environment

Let's assume that the price set for the seller is:  $\mathcal{P} = \{0, 0.5, 1\}$

- Suppose that, for each time step  $t$ , the buyer's valuation  $v_t$  follows the following distribution:

$$Pr[v_t = 0] = 0.25; Pr[v_t = 0.7] = 0.6; Pr[v_t = 1.5] = 0.15$$

- We can determine a **reward distribution  $D(a)$**  for each price  $a \in \mathcal{P}$ . For example, if  $a = 0.5$ , the reward distribution is:

$$Pr[r = 1] = 0$$

$$Pr[r = 0.5] = Pr[v \geq a] = Pr[v = 0.7] + Pr[v = 1.5] = 0.75$$

$$Pr[r = 0] = Pr[v < a] = Pr[v = 0] = 0.25$$

- The **expected reward** for  $a = 0.5$  is:

$$\mu(a) = Pr_{D(a)}[r = 0] \cdot 0 + Pr_{D(a)}[r = 0.5] \cdot 0.5 + Pr_{D(a)}[r = 1] \cdot 1 = 0.375$$

# Posted Price with Deterministic Adversarial Environment

Let's assume that  $\mathcal{P} = \{0, 0.5, 1\}$  and  $T = 3$ .

- For each time step  $t$ , the valuation  $v_t$  is chosen deterministically and does not depend on the arms chosen by the seller in the previous steps
- **Adversarial setting**  $\rightarrow$  Buyer knows the pricing algorithm used by the seller
- We can think that the buyer determines its valuations before round 1.

Step	1	2	3
$\mathbf{v}$	0.5	0	1

- For each price  $a$  the rewards are deterministic. The **rewards table** is chosen before the game starts

Step	1	2	3
Posted Price = 0	0	0	0
Posted Price = 0.5	0.5	0	0.5
Posted Price = 1	0	0	1

# Performance Evaluation

How do we argue whether an algorithm is doing a good job?



# Performance Evaluation

How do we argue whether an algorithm is doing a good job?

The problem is that some problem instances inherently allow higher rewards.

## How do we argue whether an algorithm is doing a good job?

The problem is that some problem instances inherently allow higher rewards.

Consider the stochastic setting in the slide 18 with  $T = 3$ .

- The Seller always chooses the price 0.5.
- Different problem instances:
  - The buyer always chooses a valuation equal to 0  $\rightarrow$  Cumulated revenue is 0;
  - The buyer always chooses a valuation equal to 0.7  $\rightarrow$  Cumulated revenue is 1.5
  - ...

# Performance Evaluation

How do we argue whether an algorithm is doing a good job?

The problem is that some problem instances inherently allow higher rewards.

A standard approach:

- Compare the cumulative reward (or cost) to the cumulative reward obtained by always playing the optimal arm  $a^*$ , called the **best-arm benchmark**
- We analyze the **regret**  $R(t)$  that measures the difference between the cumulative reward obtained by the bandit algorithm and the best-arm benchmark.
- Let's note  $r_\tau(a^*)$  the reward obtained at step  $\tau$  with the optimal arm  $a^*$ , and  $r_\tau(a_\tau)$  the reward obtained at step  $\tau$  with the arm played by the bandit algorithm at the same step  $a_\tau$

$$R(t) = \sum_{i=\tau}^t r_\tau(a^*) - \sum_{\tau=1}^t r_\tau(a_\tau)$$

# The Regret

- The regret at round  $t$  depends on all the arms chosen up to round  $t$
- One of the core questions in the study of Bandits is to understand the growth rate of the regret as  $t$  grows
- Our goal: the difference between the cumulative reward obtained by the bandit algorithm and the best-arm benchmark decreases over time, and goes to zero as  $t$  increases.
- Formally, A good learner achieves sub-linear regret. This means that  $\lim_{t \rightarrow \infty} \frac{R(t)}{t} = 0$ . Example of sub-linear regret:  
 $R(t) = O(\sqrt{t}), R(t) = O(\log t)$

Note that the arms chosen by the algorithm are random quantities, as they may depend on randomness in rewards and/or in the algorithm. We will typically talk about expected regret  $\mathbb{E}[R(T)]$ .

In words: *regret measures how much the algorithm "regrets" not knowing the best arm in advance*

# Stochastic Bandits

# Stochastic Bandits: Problem Structure



- **Bandit Feedback:** The algorithm observes only the reward for the selected action, and nothing else.
- **Reward Distributions:** For each action  $a$ , there is an **unknown distribution**  $D_a$  called the reward distribution. Every time action  $a$  is chosen, the reward is sampled independently from this distribution. The **expected reward** for  $a$  is  $\mu(a)$
- **Rewards are bounded:** For simplicity, we restrict the rewards to the interval  $[0, 1]$
- **Goal:** Maximize (Minimize) the cumulated revenue (cost)  $\sum_{t=1}^T r_t$

---

## Problem Protocol: Stochastic Bandits

---

**Require:**  $K$  arms;  $T$  rounds; reward distribution  $D_a$  for each arm  $a$  (unknown).

- 1: **for**  $t \in [T]$  **do**
  - 2:     Algorithm chooses some arm  $a_t \in [K]$
  - 3:     Reward  $r_t \in [0, 1]$  is sampled independently from distribution  $D_a, a = a_t$
  - 4:     Algorithm collects  $r_t$ , and observe nothing else
  - 5: **end for**
-

# Explore-First Algorithm

**Idea:** Try each arms a fixed number of times, then choose the best one.



# Explore-First Algorithm

**Idea:** Try each arms a fixed number of times, then choose the best one.

---

## Explore-First Algorithm

---

**Require:**  $K$  arms;  $T$  rounds; length of exploration phase for each arm  $N$ .

- 1: **Exploration phase:** try each arm  $N$  times
  - 2: Select the arm  $\bar{a}$  with the highest average reward (break ties arbitrarily)
  - 3: **Exploitation phase:** play arm  $\bar{a}$  in all remaining rounds
-

# Explore-First Algorithm

**Idea:** Try each arms a fixed number of times, then choose the best one.

---

## Explore-First Algorithm

---

**Require:**  $K$  arms;  $T$  rounds; length of exploration phase for each arm  $N$ .

- 1: **Exploration phase:** try each arm  $N$  times
  - 2: Select the arm  $\bar{a}$  with the highest average reward (break ties arbitrarily)
  - 3: **Exploitation phase:** play arm  $\bar{a}$  in all remaining rounds
- 

## Theorem

*Explore-first achieves regret*

$$\mathbb{E}[R(T)] \leq T^{2/3} O(K \cdot \log T)^{1/3}$$

# Epsilon-greedy algorithm

**Problem with Explore-first:** the performance in the exploration phase may be very bad if many/most of the arms have a low reward (high cost) with respect to the optimal-arm. It is usually better to spread the exploration of the arms more uniformly over the time steps.

# Epsilon-greedy algorithm

**Problem with Explore-first:** the performance in the exploration phase may be very bad if many/most of the arms have a low reward (high cost) with respect to the optimal-arm. It is usually better to spread the exploration of the arms more uniformly over the time steps.

---

## Epsilon-Greedy Algorithm

---

**Require:**  $K$  arms,  $T$  rounds,  $\epsilon_t$ .

- 1: **for** each round  $t = 1, 2, \dots, T$  **do**
  - 2:     Toss a coin with success probability  $\epsilon_t$
  - 3:     **if** success **then:**
  - 4:         **explore:** choose an arm uniformly at random
  - 5:     **else**
  - 6:         **exploit:** choose the arm with the highest average reward so far
  - 7:     **end if**
  - 8: **end for**
-

# Epsilon-greedy algorithm: The Regret

We can derive the same regret bound as for Explore-first, but now it **holds for all rounds  $t$**

## Theorem

*Epsilon-Greed algorithm achieves regret bound*

$$\mathbb{E}[R(t)] \leq t^{2/3} O(K \cdot \log t)^{1/3}$$

# Non-Adaptive Exploration

- **Non-Adaptive Exploration**: the number of exploration rounds is fixed before the game starts
- All the algorithms showed in the previous slides are Non-Adaptive Algorithms

## Claim

For any non-adaptive algorithms we can not do better than  $\Omega(T^{2/3}K^{1/3})$

For achieving better performance we must use **Adaptive Algorithms**.

# Adaptive Exploration

- Let  $n_t(a)$  the number of times when action the action  $a$  was chosen up to the round  $t$ .  $[n_t(a)]$  is the set of time steps when the action  $a$  was chosen up to the round  $t$
- The average reward of  $a$  at step  $t$  is

$$\bar{\mu}_t(a) = \frac{1}{n_t(a)} \sum_{t \in [n_t(a)]} r_t$$

- The confidence radius of arm  $a$ :  $rad_t(a) = \sqrt{2 \log T / n_t(a)}$
- With high probability  $\bar{\mu}_t(a) - rad_t(a) \leq \mu(a) \leq \bar{\mu}_t(a) + rad_t(a)$
- Upper Confidence Bound :  $UCB_t(a) = \bar{\mu}_t(a) + rad_t(a)$ .
- UCB is an upper bound of  $\mu(a)$

# Adaptive Exploration

- Let  $n_t(a)$  the number of times when action the action  $a$  was chosen up to the round  $t$ .  $[n_t(a)]$  is the set of time steps when the action  $a$  was chosen up to the round  $t$
- The average reward of  $a$  at step  $t$  is

$$\bar{\mu}_t(a) = \frac{1}{n_t(a)} \sum_{t \in [n_t(a)]} r_t$$

- The confidence radius of arm  $a$ :  $rad_t(a) = \sqrt{2 \log T / n_t(a)}$
- With high probability  $\bar{\mu}_t(a) - rad_t(a) \leq \mu(a) \leq \bar{\mu}_t(a) + rad_t(a)$
- Upper Confidence Bound :  $UCB_t(a) = \bar{\mu}_t(a) + rad_t(a)$ .
- UCB is an upper bound of  $\mu(a)$  with high probability



# Adaptive Exploration

- Let  $n_t(a)$  the number of times when action the action  $a$  was chosen up to the round  $t$ .  $[n_t(a)]$  is the set of time steps when the action  $a$  was chosen up to the round  $t$
- The average reward of  $a$  at step  $t$  is

$$\bar{\mu}_t(a) = \frac{1}{n_t(a)} \sum_{t \in [n_t(a)]} r_t$$

- The confidence radius of arm  $a$ :  $rad_t(a) = \sqrt{2 \log T / n_t(a)}$
- With high probability  $\bar{\mu}_t(a) - rad_t(a) \leq \mu(a) \leq \bar{\mu}_t(a) + rad_t(a)$
- Upper Confidence Bound :  $UCB_t(a) = \bar{\mu}_t(a) + rad_t(a)$ .
- UCB is an upper bound of  $\mu(a)$  with high probability
- If  $n_t(a) \rightarrow \infty$ :

# Adaptive Exploration

- Let  $n_t(a)$  the number of times when action the action  $a$  was chosen up to the round  $t$ .  $[n_t(a)]$  is the set of time steps when the action  $a$  was chosen up to the round  $t$
- The average reward of  $a$  at step  $t$  is

$$\bar{\mu}_t(a) = \frac{1}{n_t(a)} \sum_{t \in [n_t(a)]} r_t$$

- The confidence radius of arm  $a$ :  $rad_t(a) = \sqrt{2 \log T / n_t(a)}$
- With high probability  $\bar{\mu}_t(a) - rad_t(a) \leq \mu(a) \leq \bar{\mu}_t(a) + rad_t(a)$
- Upper Confidence Bound :  $UCB_t(a) = \bar{\mu}_t(a) + rad_t(a)$ .
- UCB is an upper bound of  $\mu(a)$  with high probability
- If  $n_t(a) \rightarrow \infty$ :  $UCB_t(a) \rightarrow \mu(a)$

An arm  $a$  can have a large  $UCB_t(a) = \overline{\mu}_t(a) + rad_t(a)$  for two reasons:

- $\overline{\mu}_t(a)$  is large:
- $rad_t(a)$  is large:

An arm  $a$  can have a large  $UCB_t(a) = \overline{\mu}_t(a) + rad_t(a)$  for two reasons:

- $\overline{\mu}_t(a)$  is large: the arm  $a$  is likely to have a high reward
- $rad_t(a)$  is large:

An arm  $a$  can have a large  $UCB_t(a) = \overline{\mu}_t(a) + rad_t(a)$  for two reasons:

- $\overline{\mu}_t(a)$  is large: the arm  $a$  is likely to have a high reward
- $rad_t(a)$  is large: the arm  $a$  has not been explored much

The two terms in  $UCB_t(a)$  represent, resp., **exploitation and exploration**.

An arm  $a$  can have a large  $UCB_t(a) = \overline{\mu}_t(a) + rad_t(a)$  for two reasons:

- $\overline{\mu}_t(a)$  is large: the arm  $a$  is likely to have a high reward
- $rad_t(a)$  is large: the arm  $a$  has not been explored much

The two terms in  $UCB_t(a)$  represent, resp., **exploitation and exploration**.

**Idea:** [*Optimism in the face of uncertainty*] At each time step  $t$  takes the arm  $a$  that maximize  $UCB_t(a)$

---

**UCB1**


---

**Require:**  $K$  arms, horizon  $T$

- 1: **for** each round  $t = 1, \dots, T$  **do**
- 2:     Chose  $a_t$  as the arm  $a$  which maximizes  $UCB_t(a)$
- 3:     Observe the reward  $r_t$
- 4:     Update upper confidence bounds:

$$\begin{cases} UCB_{t+1}(a) = \frac{1}{n_t(a)+1} \sum_{t \in [n_t(a)] \cup t} r_t + \sqrt{\frac{2 \log T}{n_t(a)+1}} & \text{if } a = a_t, \\ UCB_{t+1}(a) = UCB_t(a) & \text{otherwise} \end{cases}$$

5: **end for**

---



---

Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. “Finite-Time Analysis of the Multiarmed Bandit Problem”. In: *Mach. Learn.* 47.2–3 (2002), 235–256. ISSN: 0885-6125. DOI: 10.1023/A:1013689704352. URL: <https://doi.org/10.1023/A:1013689704352>

## Theorem

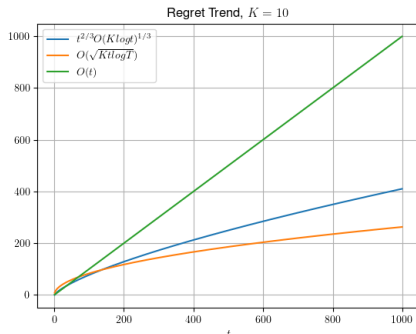
*UCB1 satisfies achieves regret bound*

$$\mathbb{E}[R(t)] = O(\sqrt{K \cdot t \log T}) \text{ for all } t \leq T$$



# Stochastic Bandits: Summary

- We use the notion of regret for evaluating the performance of a Bandits Algorithms
- Our aim is to achieve sub-linear regret
- **Non-Adaptive Algorithm:**  
 $\mathbb{E}[R(t)] \leq t^{2/3} O(K \cdot \log t)^{1/3}$
- **Adaptive-Algorithm:**  
 $\mathbb{E}[R(t)] \leq O(\sqrt{K \cdot t \cdot \log T})$




## \*Exercise

**Exercise.** Consider the Repeated Posted Price Auction described in Slide 5 with a discrete price set  $\{0.2, 0.4, 0.6, 0.8, 1.0\}$ . Specifically, for each time step  $t$  the price  $\pi_t$  belongs to  $\{0.2, 0.4, 0.6, 0.8, 1.0\}$ . Moreover, assume that the buyer samples its valuation from a fixed distribution  $D(\cdot)$ .

- 1 Suppose that you are the seller. Design a Bandit Algorithm for the Repeated Posted Price Auction
- 2 Determine the Regret of the proposed algorithm
- 3 What happens if the price set is continuous? i.e. for each time step  $t$  the price  $\pi_t$  belongs to the continuous set  $(0, 1]$

# Adversarial Bandits

# Adversarial Bandits: Problem Structure



- **Bandit Feedback:** The algorithm only observes the reward for the selected action
- **Deterministic Environment:** The environment determines the rewards deterministically, independently from the arms chosen by the learner
- **Adversarial Environment:** The environment knows the learner algorithm
- **Rewards are bounded:** For simplicity, we restrict the rewards to the interval  $[0, 1]$
- **Goal:** Maximize the cumulated revenue  $\sum_{t=1}^T r_t$

# Adversarial Environment

## The Adversarial nature of the environment is crucial

- Let's back to the Repeated Posted Price Auction
- Suppose that the Seller algorithm is deterministic:
  - The minimum price is  $p$
  - The Seller adapts the posted price by observing the rewards  $\rightarrow$  if the buyer always have the same valuation, then the posted price converges to the buyer valuation
- The Buyer knows the Seller's algorithm

## The Adversarial nature of the environment is crucial

- Let's back to the Repeated Posted Price Auction
- Suppose that the Seller algorithm is deterministic:
  - The minimum price is  $p$
  - The Seller adapts the posted price by observing the rewards  $\rightarrow$  if the buyer always have the same valuation, then the posted price converges to the buyer valuation
- The Buyer knows the Seller's algorithm
- The buyer does not act truthfully: it imposes its valuation equal to  $p$
- The Seller has no chance to sell the shoes at a price higher than  $p$

# Adversarial Environment

## The Adversarial nature of the environment is crucial

- Let's back to the Repeated Posted Price Auction
- Suppose that the Seller algorithm is deterministic:
  - The minimum price is  $p$
  - The Seller adapts the posted price by observing the rewards  $\rightarrow$  if the buyer always have the same valuation, then the posted price converges to the buyer valuation
- The Buyer knows the Seller's algorithm
- The buyer does not act truthfully: it imposes its valuation equal to  $p$
- The Seller has no chance to sell the shoes at a price higher than  $p$

## Claim

The learner have to use a randomised policy

# Adversarial Bandit Algorithm vs Mechanism Design

## Mechanism Design

- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility

## Adversarial Bandit Algorithm

- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility



# Adversarial Bandit Algorithm vs Mechanism Design

## Mechanism Design

- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility
- **Solution:** Determine Truthful Mechanism  $\rightarrow$  Reveal private valuation it is a dominant strategy

## Adversarial Bandit Algorithm

- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility

# Adversarial Bandit Algorithm vs Mechanism Design

## Mechanism Design


- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility
- **Solution:** Determine Truthful Mechanism → Reveal private valuation it is a dominant strategy

## Adversarial Bandit Algorithm

- **Problem:** Buyers can act strategically by misreporting their valuation to increase their utility
- **Solution:** Randomize the strategy

# Warm Up: Adversarial Bandits with Full Feedback

Let's make the seller's life simpler...



- **Full Feedback:** The algorithm observes the outcome not only for the chosen arm but for all the other arms as well

We introduce a randomized algorithm called **Hedge**

- The algorithm takes into account the reward of all arms
- Hedge maintains a weight  $w_t(a)$  for each arm  $a$
- The higher is the weight for the arm the higher is the confidence about the arm
- Hedge uses a randomized algorithm to select the arm: the higher is the confidence about the arm the higher is the probability of selecting the arm
- **Idea**: Initially all the arms have the same confidence. If an arm receives an high reward we increase our confidence about the arm. At each time step we pick the arm with the higher confidence with high probability

---

## Hedge

---

**Require:**  $K$  arms,  $T$  rounds,  $\epsilon \in [0, 1/2]$

- 1: Initialize the weights  $w_i(a) = 1$  for each arm  $a$ .
  - 2: **for**  $t \in [T]$  **do**
  - 3:     Let  $p_t(a) = \frac{w_t(a)}{\sum_{a' \in [K]} w_t(a')}$
  - 4:     Sample an arm  $a_t$  from  $p_t(\cdot)$
  - 5:     **for**  $a \in [K]$  **do**
  - 6:         Observe  $r_t(a)$
  - 7:          $w_{t+1}(a) \leftarrow w_t(a) \cdot (1 - \epsilon)^{(1-r_t(a))}$
  - 8:     **end for**
  - 9: **end for**
- 

---

Nicolò Cesa-Bianchi et al. "How to Use Expert Advice". In: *J. ACM* 44.3 (1997), 427–485. ISSN: 0004-5411. DOI: 10.1145/258128.258179. URL: <https://doi.org/10.1145/258128.258179>

## Theorem

*Assume that the reward is lower bounded, i.e.  $r_t \geq c$ , for some known  $c > 0$  and for all  $t$ . Then Hedge achieves regret:*

$$\mathbb{E}[R(T)] \leq O(2 \cdot \sqrt{T \cdot \log K})$$

# Reduction from Bandit Feedback to Full Feedback

- Use Hedge to solve a Bandit Adversarial Problem
- Sample an action according to the probability distribution defined by Hedge
- **Problem:** At each time step Hedge needs the reward of all the arms
- **Solution:** Define a fake reward for all the arms

# Reduction from Bandit Feedback to Full Feedback

---

## Reduction Schema

---

**Require:**  $K$  arms,  $T$  rounds,  $\epsilon > 0$  for Hedge

- 1: **for**  $t \in [T]$  **do**
  - 2:     Call Hedge, receive the probability distribution  $p_t$  over  $[K]$
  - 3:     **Selection rule:** use  $p_t$  to pick arm  $a_t$
  - 4:     Observe the reward  $r_t(a_t)$  of the chosen arm (*Bandit Feedback*)
  - 5:     **Fake rewards:** Define *fake rewards*  $\hat{r}_t(a)$  for all arms
  - 6:     Return *fake rewards* to Hedge
  - 7: **end for**
-



Exp3 completes the Reduction Schema:

- **Selection Rule:** (Explore and randomize) With probability  $\gamma \in [0, 1/2)$  pick an arm uniformly at random, otherwise draw an arm from  $p_t$
- **Fake rewards:**

$$\hat{r}_t(a) = \begin{cases} 1 - \frac{1-r_t(a)}{p_t(a)}, & \text{if } a = a_t, \\ 1, & \text{otherwise} \end{cases}$$

---

Peter Auer et al. "The Nonstochastic Multiarmed Bandit Problem". In: *SIAM J. Comput.* 32.1 (2003), 48–77. ISSN: 0097-5397. DOI: 10.1137/S0097539701398375.  
 URL: <https://doi.org/10.1137/S0097539701398375>

## Exp3

**Require:**  $K$  arms,  $T$  rounds,  $\epsilon \in [0, 1/2]$  for Hedge,  $\gamma \in [0, 1/2)$

- 1: **for**  $t \in [T]$  **do**
- 2:     Call Hedge, receive the probability distribution  $p_t$  over  $[K]$
- 3:     With prob.  $\gamma$  pick an arm  $a_t$  uniformly at random; otherwise draw an arm  $a_t$  from  $p_t$
- 4:     Observe the reward  $r_t(a_t)$  of the chosen arm
- 5:     Compute the fake rewards for all arms  $a \in [K]$ :

$$\hat{r}_t(a) = \begin{cases} 1 - \frac{1-r_t(a)}{p_t(a)}, & \text{if } a = a_t, \\ 1, & \text{otherwise} \end{cases}$$

- 6:     Return *fake rewards* to Hedge
- 7: **end for**

## Theorem

*Assume that the reward is lower bounded, then Exp3 achieves regret:*

$$\mathbb{E}[R(T)] \leq O(2\sqrt{T \cdot K \cdot \log K})$$

# Exp3 Parameters

How to impose parameters  $\epsilon$  and  $\gamma$ ? Maths can help!

- Remark  $r_t \in [c, 1], \forall t$
- $\gamma \in [0, \frac{1}{2T})$
- $\epsilon = \sqrt{\frac{(1-\gamma) \cdot \log K}{3KT}}$

This choice of parameters guarantees the regret showed in the last slide!

# Adversarial Bandits: Summary

- The learner has to randomise its strategy;
- We reduce our Bandit Problem to a Full-Feedback Problem
- Exp3:

$$\mathbb{E}[R(T)] \leq O(2\sqrt{T \cdot K \cdot \log K})$$

## \*Exercise

**Exercise.** Consider the Repeated Posted Price Auction described in Slide 5 with a discrete price set  $\{0.2, 0.4, 0.6, 0.8, 1.0\}$ . Specifically, for each time step  $t$  the price  $\pi_t$  belongs to  $\{0.2, 0.4, 0.6, 0.8, 1.0\}$ . Moreover, assume that the buyer determines its valuation for all weeks before the game starts, the buyer knows the algorithm used by the seller and the buyer's valuation is at least 0.2

- 1 Suppose that you are the seller. Design a Bandit Algorithm for the Repeated Posted Price Auction
- 2 Determine the Regret of the proposed algorithm
- 3 What happens if the price set is continuous? i.e. for each time step  $t$  the price  $\pi_t$  belongs to the continuous set  $(0, 1]$

## References

## Reference Books:

- Tor Lattimore and Csaba Szepesvári. *Bandit Algorithms*. Cambridge University Press, 2019,  
<https://tor-lattimore.com/downloads/book/book.pdf>
- Aleksandrs Slivkins. *Introduction to Multi-Armed Bandits*. Foundations and Trends in Machine Learning, 2019,  
<https://arxiv.org/abs/1904.07272>